

▼ Missing data imputation in Probit regression

In this programming assignment, we will implement a Markov chain Monte Carlo (MCMC) algorithm for binary regression with missing features. Refer to the slides for lecture 23.

Note: For submitting your solutions, do the following:

1. Add link to this colab notebook in the last cell of this notebook.
2. **File -> Save -> Print -> Save as PDF** and upload the pdf to gradescope.

(No work) **Step 1:** Generate data: we have provided code to randomly generate data, including the ground truth regression coefficient vector, feature matrix, and binary outcomes. We also provide you the `truncrandn` function; it generates random numbers from the truncated normal distribution. It has the following input-output format; you can input scalars but the function is implemented in a way such that you can input vectors and generate multiple random numbers at the same time.

We also give you ways to generate from the inverse Wishart distribution; code is provided for that part.

```

1 import scipy.stats
2 import numpy as np
3 import numpy.random as npr
4 from numpy.linalg import inv, cholesky
5 from scipy.stats import chi2
6 import matplotlib.pyplot as plt
7 import math
8
9 def invwishartrand(nu, phi):
10     return inv(wishartrand(nu, inv(phi)))
11
12 def wishartrand(nu, phi):
13     dim = phi.shape[0]
14     chol = cholesky(phi)
15     foo = np.zeros((dim,dim))
16
17     for i in range(dim):
18         for j in range(i+1):
19             if i == j:
20                 foo[i,j] = np.sqrt(chi2.rvs(nu-(i+1)+1))
21             else:
22                 foo[i,j] = npr.normal(0,1)
23     return np.dot(chol, np.dot(foo, np.dot(foo.T, chol.T)))
24
25 def truncrandn(mu,sigma,indic):
26

```

```

27 L = len(indic)
28 tempno = np.zeros((L,))
29 idx1 = np.where(indic==1)[0]
30 idx0 = np.where(indic==0)[0]
31
32 # generate the positive side
33 resid = scipy.stats.norm.cdf(0.0,loc = mu[idx1], scale = sigma[idx1])
34 resid1 = resid + np.random.uniform(size=(len(idx1),)) * (1.0 - resid)
35 tempno[idx1] = scipy.stats.norm.ppf(resid1,loc = mu[idx1], scale = sigma[idx1])
36
37 # generate the negative side
38 resid = scipy.stats.norm.cdf(0.0,loc = mu[idx0], scale = sigma[idx0])
39 resid0 = np.random.uniform(size=(len(idx0),)) * resid
40 tempno[idx0] = scipy.stats.norm.ppf(resid0,loc = mu[idx0], scale = sigma[idx0])
41
42 idxinf = np.where(np.isinf(tempno))[0]
43 if len(idxinf) > 0:
44     # flip to the negative side to sample!!!
45     resid = scipy.stats.norm.cdf(0.0,loc = -mu[idxinf], scale = sigma[idxinf])
46     resid_inf = np.random.uniform(size=(len(idxinf),)) * resid
47     tempno[idxinf] = 2.0 * mu[idxinf] - \
48         scipy.stats.norm.ppf(resid_inf,loc = mu[idxinf], scale = sigma[idxinf])
49
50 return tempno
51
52
53 1 # generate data
54 2 np.random.seed(0)
55 3 N = 200
56 4 P = 10
57 5 pct_miss = .2
58 6
59 7 X = np.random.normal(size=(N,P))
60 8 beta = .2 * np.random.normal(size=(P,))
61 9 y = np.random.binomial(1,1.0/(1+np.exp(-np.dot(X,beta))))
62 10 X[np.random.binomial(1,pct_miss,size=(N,P)) == 1] = np.nan

```

Step 2: Baseline. Use what you have implemented in PA1 to run logistic regression on this data. For a missing feature in an observation, do an imputation by taking the average of this feature's values across all observations where it is not missing. Implement it as logreg mixx function, with two inputs: X and y , the former is the feature matrix with missing entries in it while the latter is the binary-valued vector of observations.

```

1 z = lambda y, A, x: np.exp((2*y - 1) * (A @ x))
2 f = lambda y, A, x: np.sum(np.log(1 + np.exp(-(2*y - 1) * (A @ x))))
3 fp = lambda y, A, x: -A.T @ ((2*y - 1) / (1 + z(y, A, x)))
4 fpp = lambda y, A, x: A.T @ (np.diag(1 / (2 + z(y, A, x) + 1/z(y, A, x)))) @ A
5

```

```

6 def logreg_misx(Xin,y):
7
8     # l2 regularized logistic regression
9     N,P = Xin.shape
10
11     # # impute with simple averages
12     X = np.copy(Xin)
13     for ii in range(N):
14         for jj in range(P):
15             if np.isnan(Xin[ii,jj]):
16                 X[ii,jj] = np.nanmean(Xin[:,jj])
17
18     """
19     Add your code here
20     """
21     x = np.zeros(P)
22     A = X
23     it = 0
24     ss = 10
25     tol = 1e-15
26     change = math.inf
27     alpha, beta = 1, 1/2
28     while it < 10000 and change > tol:
29         grad = fp(y, A, x)
30         while f(y, A, x - ss * grad) > f(y, A, x) - (alpha/2)*ss*(np.linalg.norm(grad
31             ss = beta * ss
32         x_change = - ss * fp(y, A, x)
33         new_x = x + x_change
34         old_obj = f(y, A, x)
35         change = abs(f(y, A, new_x) - old_obj) / abs(old_obj)
36         it += 1
37         x = new_x
38     return x
39

```

Step 3: Implement the MCMC algorithm. We have provided all other parts of the code; all you need to do is to implement sampling steps for z_i , β , and $x_i^{M_i}$ for all i .

```

1 def proreg_misx_b(X_input,y):
2
3     # the main function
4     (N,P) = X_input.shape
5
6     # some algo parameters
7     T = 1000
8     burnin = int(T/2)
9
10     mu_beta = np.zeros((P,))
11     Sig_beta = np.eye(P)

```

```

12     Pres_beta = np.linalg.inv(Sig_beta)
13     kappa_0 = 1.0
14     nu_0 = 1.0
15     mu_0 = np.zeros((P,))
16     Psi_0 = np.eye(P)
17
18     # allocate variables
19     mask = np.double(~np.isnan(X_input)) # this is a pre-calculation
20     X = np.zeros((N,P,T))
21     z = np.zeros((N,T))
22     beta = np.zeros((P,T))
23     mu = np.zeros((P,T))
24     Sig = np.zeros((P,P,T))
25
26     # initialize
27     # X initialization: simply average the unobserved X entries
28     for ii in range(N):
29         for jj in range(P):
30             if np.isnan(X_input[ii,jj]):
31                 X[ii,jj,0] = np.nanmean(X_input[:,jj]) #.05 * np.random.normal()
32             else:
33                 X[ii,jj,0] = X_input[ii,jj]
34
35     # beta initialization: initialize with log reg
36     beta[:,0] = logreg_misx(X_input,y)
37     # beta[:,0] = np.zeros((P,))
38     z[:,0] = truncrandn(np.dot(X[:, :, 0], beta[:, 0]), np.ones((N,)), y)
39
40     # mu and Sigma calculated according to imputations
41     mu[:,0] = np.mean(X[:, :, 0], axis=0)
42     Sig[:, :, 0] = np.dot(X[:, :, 0].T, X[:, :, 0]) / np.double(N) + .1 * np.eye(P)
43
44     for tt in range(1,T):
45
46         if tt%100 == 0:
47             print('iteration ', tt)
48
49         # sample z
50         """
51         Add your code here
52         """
53         z[:,tt] = truncrandn(np.dot(X[:, :, tt-1], beta[:, tt-1]), np.ones((N,)), y)
54         znow = np.copy(z[:,tt])
55         Xnow = np.copy(X[:, :, tt-1])
56         # sample beta
57         """
58         Add your code here
59         """
60         Pres_n = np.dot(Xnow.T, Xnow) + Pres_beta
61         Sig_n = np.linalg.inv(Pres_n)
62         mu_n = np.dot(Sig_n, np.dot(Pres_beta, mu_beta) + np.dot(Xnow.T, znow))

```

```

63     beta[:,tt] = np.random.multivariate_normal(mu_n,Sig_n)
64     Pres_beta = Pres_n
65     Sig_beta = Sig_n
66     mu_beta = mu_n
67
68     # now sample the missing X values
69     """
70     Add your code here
71     """
72     for ii in range(N):
73         mis_idx = np.where(mask[ii,] == 0)[0]
74         if len(mis_idx) > 0:
75             obs_idx = np.where(mask[ii,] > 0)[0]
76             mum = mu[:,tt-1][mis_idx]
77             muo = mu[:,tt-1][obs_idx]
78             Sigmm = Sig[:, :, tt-1][mis_idx, :][:, mis_idx]
79             Sigmo = Sig[:, :, tt-1][mis_idx, :][:, obs_idx]
80             Sigom = Sig[:, :, tt-1][obs_idx, :][:, mis_idx]
81             Sigoo = Sig[:, :, tt-1][obs_idx, :][:, obs_idx]
82             mubar = mum + np.dot(np.dot(Sigmo, inv(Sigoo)), (Xnow[ii,obs_idx] - muo))
83             Sigbar = Sigmm - np.dot(np.dot(Sigmo, inv(Sigoo)), Sigom)
84
85             zbar = xnow[ii] - np.dot(Xnow[ii, obs_idx].T, beta[obs_idx,tt])
86             Sigprime = inv(inv(Sigbar) + np.dot(beta[mis_idx, tt], beta[mis_idx,tt].T))
87             muprime = np.dot(Sigprime, np.dot(inv(Sigbar), mubar) + np.dot(zbar, beta[mis_idx,tt]))
88             X[ii,mis_idx,tt] = np.random.multivariate_normal(muprime, Sigprime)
89             X[ii, obs_idx, tt] = X_input[ii, obs_idx]
90
91
92     # now sample mu and Sig values
93     # first calculate sufficient stats
94     Xnow = X[:, :, tt]
95     xbar = np.mean(Xnow,axis=0)
96     C = np.zeros((P,P))
97     for ii in range(N):
98         C += np.outer(Xnow[ii,] - xbar,Xnow[ii,] - xbar)
99     # then use them to get new IW params
100     nu_t = nu_0 + N
101     Sig_t = Psi_0 + C + kappa_0 * N / nu_t * np.outer(xbar-mu_0,xbar-mu_0)
102     Sig[:, :, tt] = invwishartrand(nu_t,Sig_t)
103     # now use the newly sampled Sig to sample new mu
104     kappa_t = kappa_0 + N
105     mu_t = (kappa_0 * mu_0 + N * xbar) / kappa_t
106     mu[:,tt] = np.random.multivariate_normal(mu_t,Sig[:, :, tt] / kappa_t)
107
108     # run the freakin' loop on and on... till convergence!
109     betaout = np.mean(beta[:,burnin:],axis=1)
110     Xout = np.mean(X[:, :, burnin:],axis=2)
111     muout = np.mean(mu[:,burnin:],axis=1)
112     Sigout = np.mean(Sig[:, :, burnin:],axis=2)

```

```

113
114     return betaout,Xout,muout,Sigout

```

Step 4a): Investigate. Run the MCMC algorithm for a total of 1000 iterations with the first 500 as burn-in, and set your own hyperparameter values. Obtain an estimated regression coefficient vector and compare it to the ground truth and calculate the l_2 -norm of the difference. Do the same for the logistic regression baseline; what do you observe?

```

1 betahat_b = proreg_misx_b(X,y)[0]
2 betahat_o = logreg_misx(X,y)
3
4 print('MCMC probit error = ',np.linalg.norm(betahat_b-beta) / np.linalg.norm(beta))
5 print('Logistic regression error = ',np.linalg.norm(betahat_o-beta) / np.linalg.no

iteration    100
iteration    200
iteration    300
iteration    400
iteration    500
iteration    600
iteration    700
iteration    800
iteration    900
MCMC probit error =  0.32077982842541003
Logistic regression error =  0.558257587166397

```

Answer the questions and discuss your findings here

The MCMC probit error is lower than the logistic regression error

Step 4b): Plot the values of the randomly generated samples of latent variables over iterations. Pick three entries in β and discuss whether you think the number of iterations and burnin (1000 and 500) are enough.

```

1 def proreg_misx_b_alt(X_input,y):
2
3     # the main function
4     (N,P) = X_input.shape
5
6     # some algo parameters
7     T = 1000
8     burnin = int(T/2)
9
10    mu_beta = np.zeros((P,))
11    Sig_beta = np.eye(P)
12    Pres_beta = np.linalg.inv(Sig_beta)
13    kappa_0 = 1.0
14    nu_0 = 1.0

```

```

15     mu_0 = np.zeros((P,))
16     Psi_0 = np.eye(P)
17
18     # allocate variables
19     mask = np.double(~np.isnan(X_input)) # this is a pre-calculation
20     X = np.zeros((N,P,T))
21     z = np.zeros((N,T))
22     beta = np.zeros((P,T))
23     mu = np.zeros((P,T))
24     Sig = np.zeros((P,P,T))
25
26     # initialize
27     # X initialization: simply average the unobserved X entries
28     for ii in range(N):
29         for jj in range(P):
30             if np.isnan(X_input[ii,jj]):
31                 X[ii,jj,0] = np.nanmean(X_input[:,jj]) #.05 * np.random.normal()
32             else:
33                 X[ii,jj,0] = X_input[ii,jj]
34
35     # beta initialization: initialize with log reg
36     beta[:,0] = logreg_misx(X_input,y)
37     # beta[:,0] = np.zeros((P,))
38     z[:,0] = truncrandn(np.dot(X[:, :, 0], beta[:, 0]), np.ones((N,)), y)
39
40     # mu and Sigma calculated according to imputations
41     mu[:,0] = np.mean(X[:, :, 0], axis=0)
42     Sig[:, :, 0] = np.dot(X[:, :, 0].T, X[:, :, 0]) / np.double(N) + .1 * np.eye(P)
43
44     for tt in range(1,T):
45
46         if tt%100 == 0:
47             print('iteration ', tt)
48
49         # sample z
50         """
51         Add your code here
52         """
53         z[:,tt] = truncrandn(np.dot(X[:, :, tt-1], beta[:, tt-1]), np.ones((N,)), y)
54         znow = np.copy(z[:,tt])
55         Xnow = np.copy(X[:, :, tt-1])
56         # sample beta
57         """
58         Add your code here
59         """
60         Pres_n = np.dot(Xnow.T, Xnow) + Pres_beta
61         Sig_n = np.linalg.inv(Pres_n)
62         mu_n = np.dot(Sig_n, np.dot(Pres_beta, mu_beta) + np.dot(Xnow.T, znow))
63         beta[:,tt] = np.random.multivariate_normal(mu_n, Sig_n)
64         Pres_beta = Pres_n
65         Sig_beta = Sig_n

```

```

66     mu_beta = mu_n
67
68     # now sample the missing X values
69     """
70     Add your code here
71     """
72     for ii in range(N):
73         mis_idx = np.where(mask[ii,] == 0)[0]
74         if len(mis_idx) > 0:
75             obs_idx = np.where(mask[ii,] > 0)[0]
76             mum = mu[:,tt-1][mis_idx]
77             muo = mu[:,tt-1][obs_idx]
78             Sigmm = Sig[:, :, tt-1][mis_idx, :][:, mis_idx]
79             Sigmo = Sig[:, :, tt-1][mis_idx, :][:, obs_idx]
80             Sigom = Sig[:, :, tt-1][obs_idx, :][:, mis_idx]
81             Sigoo = Sig[:, :, tt-1][obs_idx, :][:, obs_idx]
82             mubar = mum + np.dot(np.dot(Sigmo, inv(Sigoo)), (X[ii, obs_idx, tt] - mu
83             Sigbar = Sigmm - np.dot(np.dot(Sigmo, inv(Sigoo)), Sigom)
84
85             #missed implementation
86             zbar = z[ii, tt] - np.dot(X[ii, obs_idx, tt].T, beta[obs_idx, tt])
87             Sigprime = inv(inv(Sigbar) + np.dot(beta[mis_idx, tt], beta[mis_idx, tt].T))
88             muprime = np.dot(Sigprime, np.dot(inv(Sigbar), mubar) + np.dot(zbar, beta[mis_idx, tt]))
89             ##
90             X[ii, mis_idx, tt] = np.random.multivariate_normal(muprime, Sigprime)
91             X[ii, obs_idx, tt] = X_input[ii, obs_idx, tt]
92
93
94     # now sample mu and Sig values
95     # first calculate sufficient stats
96     Xnow = X[:, :, tt]
97     xbar = np.mean(Xnow, axis=0)
98     C = np.zeros((P, P))
99     for ii in range(N):
100         C += np.outer(Xnow[ii,] - xbar, Xnow[ii,] - xbar)
101     # then use them to get new IW params
102     nu_t = nu_0 + N
103     Sig_t = Psi_0 + C + kappa_0 * N / nu_t * np.outer(xbar - mu_0, xbar - mu_0)
104     Sig[:, :, tt] = invwishartrand(nu_t, Sig_t)
105     # now use the newly sampled Sig to sample new mu
106     kappa_t = kappa_0 + N
107     mu_t = (kappa_0 * mu_0 + N * xbar) / kappa_t
108     mu[:, tt] = np.random.multivariate_normal(mu_t, Sig[:, :, tt] / kappa_t)
109
110     # run the freakin' loop on and on... till convergence!
111     betaout = np.mean(beta[:, burnin:], axis=1)
112     Xout = np.mean(X[:, :, burnin:], axis=2)
113     muout = np.mean(mu[:, burnin:], axis=1)
114     Sigout = np.mean(Sig[:, :, burnin:], axis=2)
115

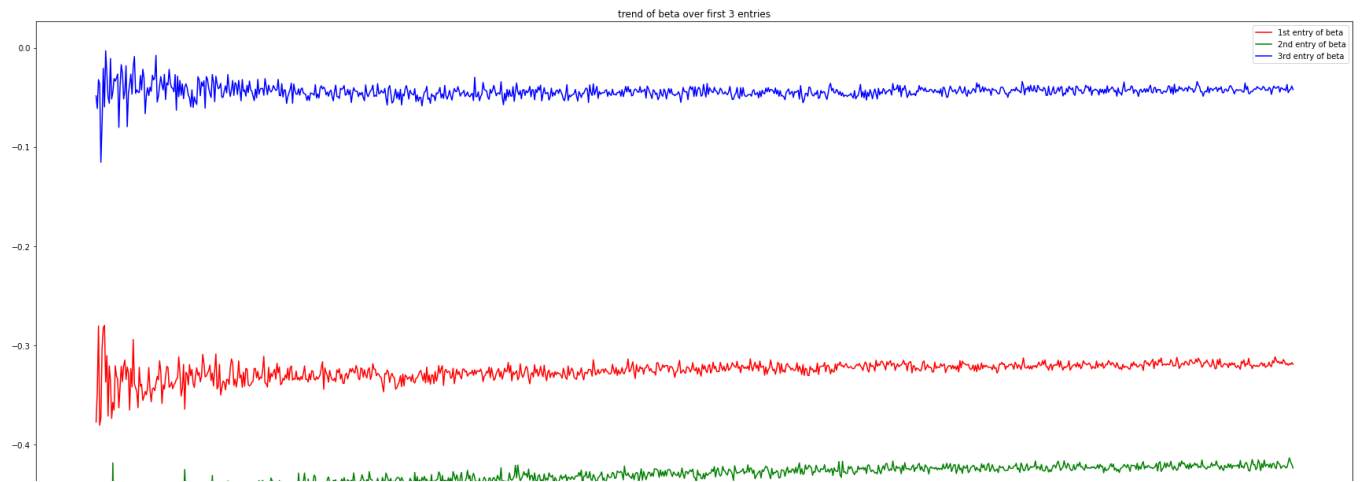
```



```
116     return beta, X, mu, Sig
117 betaout, Xout, muout, Sigout = proreg_misx_b_alt(X,y)

    iteration 100
    iteration 200
    iteration 300
    iteration 400
    iteration 500
    iteration 600
    iteration 700
    iteration 800
    iteration 900

1 """
2 Add your code here
3 """
4 beta1 = betaout[0]
5 beta2 = betaout[1]
6 beta3 = betaout[2]
7
8 #fig, axs = plt.subplots(1, figsize = (30, 30))
9
10 plt.figure(figsize = (30, 15))
11 plt.plot(beta1,'r', label = '1st entry of beta')
12 plt.plot(beta2,'g', label = '2nd entry of beta')
13 plt.plot(beta3,'b', label = '3rd entry of beta')
14 plt.title("trend of beta over first 3 entries")
15 plt.legend()
16 plt.show()
```



*Answer the questions and discuss your findings here

The individual beta are converging, which implies that 1000 iterations and burn of 500 is enough


Step 4c): Investigate the posterior distributions. Again, pick some entries in β and plot histograms of their density. Compare these with the prior and the ground-truth values; what do you observe?

```

1 """
2 Add your code here
3 """
4 print(f"first entry of ground truth beta: {beta[0]}")
5 print(f"second entry of ground truth beta: {beta[1]}")
6 print(f"third entry of ground truth beta: {beta[2]}")
7
8 fig, axs = plt.subplots(3, figsize = (30, 15))
9 axs[0].hist(beta1[500:])
10 axs[1].hist(beta2[500:])
11 axs[2].hist(beta3[500:])
12 plt.show()
13

```

first entry of ground truth beta: -0.3065842106860256
second entry of ground truth beta: -0.3423940328188443
third entry of ground truth beta: 0.009227011791136986



Answer the questions and discuss your findings here

The posterior distribution is distributed differently from the prior, but is somewhat close to the groundtruth beta

Add Colab link here:

https://colab.research.google.com/drive/1fhEHc_3Uy9tMh42x4XJ2-YP94MqyL3B8?authuser=1#scrollTo=hauUvRmh9UC

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 4:56 PM

